

## Safety & Ethics Considerations

### 1. Explicit Gender Always Overrides Inference

If a user provides gender, no inference occurs. The system never overwrites explicit self-identification.

### 2. No Storage of Sensitive Data

Photos, facial embeddings, and classifier outputs are never stored. All inference is transient.

### 3. Confidence Transparency

All inferred results include a confidence score and per-signal attribution.

### 4. Ethical Fallback

If confidence is below threshold, the system returns "Unknown" instead of guessing.

### 5. Bias Mitigation

The model combines multiple signals to reduce over-reliance on any single source.

### 6. User-Controlled Overrides

User or admin overrides persist and prevent re-inference unless explicitly cleared.

### 7. Auditability

The system outputs complete attribution, enabling debugging, transparency, and accountability.